

Advanced Machine Learning

Lecture 4: Combining Models

Gwenn Englebienne
Mannes Poel

University of Twente

Introduction

Bias-Variance Decomposition

Bagging and Boosting

Bagging

Boosting

Tree-based models

Classification trees

Random Forests

Introduction

Bias-Variance Decomposition

Bagging and Boosting


Bagging

Boosting

Tree-based models


Classification trees

Random Forests




What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ Committee: Combine the output of multiple classifiers
 - ▶ For example, average the outputs
 - ▶ Alternatively, create a “meta-classifier”




What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ **Committee: Combine the output of multiple classifiers**
 - ▶ For example, average the outputs
 - ▶ Alternatively, create a “meta-classifier”




What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ **Committee: Combine the output of multiple classifiers**
 - ▶ For example, average the outputs
 - ▶ Alternatively, create a “meta-classifier”




What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ **Committee: Combine the output of multiple classifiers**
 - ▶ For example, average the outputs
 - ▶ Alternatively, create a “meta-classifier”



What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ **Committee: Combine the output of multiple classifiers**
 - ▶ For example, average the outputs
 - ▶ **Alternatively, create a “meta-classifier”**



What are committees?

- ▶ Traditional approach: train a classifier to predict a class
- ▶ **Committee: Combine the output of multiple classifiers**
 - ▶ For example, average the outputs
 - ▶ **Alternatively, create a “meta-classifier”**

Consider M regression models $y_m(\mathbf{x})$, $1 \leq m \leq M$ predicting $h(\mathbf{x})$. Each individual prediction error is

$$\epsilon_m(\mathbf{x}) = h(\mathbf{x}) - y_m(\mathbf{x}) ,$$

with an averaging committee:

$$y(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M y_m(\mathbf{x})$$

Why committees?

The expected sum-squared errors are:

Individual model (average)	Committee
$E_{AV} = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{\mathbf{x}}[\epsilon_m^2(\mathbf{x})]$	$E_{COM} = \mathbb{E}_{\mathbf{x}} \left[\left(\frac{1}{M} \sum_{m=1}^M \epsilon_m(\mathbf{x}) \right)^2 \right]$

so that, if the errors are uncorrelated, we get

$$E_{COM} = \frac{1}{M^2} \mathbb{E}_{\mathbf{x}} \left[\sum_{m=1}^M \epsilon_m^2(\mathbf{x}) + 2 \sum_{m \neq n} \epsilon_m(\mathbf{x}) \epsilon_n(\mathbf{x}) \right] = \frac{1}{M} E_{AV}$$

Why committees?

The expected sum-squared errors are:

Individual model (average)

Committee

$$E_{AV} = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{\mathbf{x}}[\epsilon_m^2(\mathbf{x})] \quad E_{COM} = \mathbb{E}_{\mathbf{x}} \left[\left(\frac{1}{M} \sum_{m=1}^M \epsilon_m(\mathbf{x}) \right)^2 \right]$$

so that, if the errors are uncorrelated, we get

$$E_{COM} = \frac{1}{M^2} \mathbb{E}_{\mathbf{x}} \left[\sum_{m=1}^M \epsilon_m^2(\mathbf{x}) + \underbrace{2 \sum_{m \neq n} \epsilon_m(\mathbf{x}) \epsilon_n(\mathbf{x})}_{=0} \right] = \frac{1}{M} E_{AV}$$

In theory, committees can vastly reduce the expected error of individual classifiers

- ▶ Make the expected error arbitrarily small by increasing M
- ▶ In practice, the classifiers are highly correlated
 - ▶ The error reduction is then small
- ▶ But: it can be shown that

$$E_{AV} \geq E_{COM}$$

- ▶ We can improve the performance of committees by decreasing the correlation between the classifiers

Consider multiple training data sets $D = \{(\mathbf{x}_n, h(\mathbf{x}_n))\}$ of fixed size
Taking the expected squared loss of a model, we can decompose:

$$\mathbb{E}_D[(y_D(\mathbf{x}) - \hat{t}(\mathbf{x}))^2] = \underbrace{(\mathbb{E}_D[y_D(\mathbf{x})] - \hat{t}(\mathbf{x}))^2}_{\text{bias}^2} + \underbrace{\mathbb{E}_D[(y_D(\mathbf{x}) - \mathbb{E}_D[y_D(\mathbf{x})])^2]}_{\text{variance}}$$

Interpretation:

- ▶ The bias captures how well the model *can* perform.
Flexible models will have low bias.
- ▶ The variance captures how much the end model depends on the specific dataset.
Flexible models will have high variance.

Consider multiple training data sets $D = \{(\mathbf{x}_n, h(\mathbf{x}_n))\}$ of fixed size
Taking the expected squared loss of a model, we can decompose:

$$\mathbb{E}_D[(y_D(\mathbf{x}) - \hat{t}(\mathbf{x}))^2] = \underbrace{(\mathbb{E}_D[y_D(\mathbf{x})] - \hat{t}(\mathbf{x}))^2}_{\text{bias}^2} + \underbrace{\mathbb{E}_D[(y_D(\mathbf{x}) - \mathbb{E}_D[y_D(\mathbf{x})])^2]}_{\text{variance}}$$

Interpretation:

- ▶ The bias captures how well the model *can* perform.
Flexible models will have low bias.
- ▶ The variance captures how much the end model depends on the specific dataset.
Flexible models will have high variance.

Bias-Variance decomposition:

- ▶ Gives us insight into how a particular model generalises
 - ▶ High bias-low variance models do not learn from the data
 - ▶ Low bias-high variance models overfit on the training data
 - ▶ Optimal model flexibility (e.g., regularisation):
good bias–variance trade-off.
- ▶ Has little practical value: single training dataset
- ▶ Provides insight into why committees are useful

Optimal ensemble learning

For best ensemble performance, we want the base learners to be
flexible enough and as diverse as possible

Bias-Variance decomposition:


- ▶ Gives us insight into how a particular model generalises
 - ▶ High bias-low variance models do not learn from the data
 - ▶ Low bias-high variance models overfit on the training data
 - ▶ Optimal model flexibility (e.g., regularisation): good bias-variance trade-off.
- ▶ Has little practical value: single training dataset
- ▶ Provides insight into why committees are useful

Optimal ensemble learning


For best ensemble performance, we want the base learners to be
flexible enough and **as diverse as possible**

Bias-variance: an example

UNIVERSITY OF TWENTE.
HMI.



Bias-variance: an example



Introduction

Bias-Variance Decomposition

Bagging and Boosting

Bagging

Boosting

Tree-based models

Classification trees

Random Forests

Where do we get the base learners?

1. Single type of classifiers:
Homogeneous learners
2. Multiple types of classifiers:
Heterogeneous learners

Diversity in homogeneous learners?

- ▶ Subsample the training data
- ▶ Add randomness to the learning algorithm
- ▶ Manipulate attributes or outputs

Where do we get the base learners?

1. Single type of classifiers:

Homogeneous learners

2. Multiple types of classifiers:


Heterogeneous learners

Diversity in homogeneous learners?


- ▶ Subsample the training data
- ▶ Add randomness to the learning algorithm
- ▶ Manipulate attributes or outputs

We rarely have infinitely large training datasets. . .

- ▶ Or infinitely many. . .
- ▶ Using bootstrapping, we can create new datasets
- ▶ The correlation between datasets is then known and kept small
- ▶ **Bootstrap aggregation:**
Simply average the outcomes of classifiers trained on different bootstrap datasets



Sample N points at random from the data **with replacement**




The probability for not picking a data point is

$$p(\neg b) = (1 - 1/N)^N \approx 0.368 \quad (1)$$


The expected number of used data points is therefore

$$p(b) = 1 - p(\neg b) \approx 0.632N \quad (2)$$



In this example:

- ▶ A polynomial was fitted to 10 noisy training points (red)
- ▶ 1000 polynomials were fitted to bootstrap sets from the same 10 datapoints and averaged (blue line)





In this example:


- ▶ A polynomial was fitted to 10 noisy training points (red)
- ▶ 1000 polynomials were fitted to bootstrap sets from the same 10 datapoints and averaged (blue line)

Bagging

- ▶ Improves results with high-variance models
- ▶ No independent datasets (\Rightarrow small improvements)
- ▶ Cannot help with high bias models







Weak learner Learner that performs better than random

Strong learner Learner with accuracy $1 - \epsilon$, where ϵ is arbitrarily small

[Shapire 1990]: Weak learners in the same class as strong learners

Boosting

A technique to combine weak learners to form a strong learner


Weak learner Learner that performs better than random

Strong learner Learner with accuracy $1 - \epsilon$, where ϵ is arbitrarily small

[Shapire 1990]: Weak learners in the same class as strong learners

Boosting

A technique to combine weak learners to form a strong learner



Adaptive boosting:

- ▶ Assign each training datapoint a weight
- ▶ Iterate:
 - ▶ Train a classifier based on the weighted training data
 - ▶ Assign this classifier a weight based on how well it performs
 - ▶ Update the datapoints' weights based on how many classifiers classify it correctly

1. Set $w_n^{(1)} = \frac{1}{N}$

2. For $m = 1, \dots, M$

2.1 Fit $y_m(\mathbf{x})$ by minimising $E_m = \sum_{n \in \mathcal{M}_m} w_n^{(m)}$ (Total weight of missclassified points)

2.2 Evaluate

$$\epsilon_m = \frac{\sum_{n \in \mathcal{M}_m} w_n^{(m)}}{\sum_n w_n^{(m)}}, \quad \text{(Ratio missclassified for that classifier)}$$

$$\text{set } \alpha_m = \log \frac{1 - \epsilon_m}{\epsilon_m}$$

2.3 Update the weights

$$w_n^{(m+1)} = \begin{cases} w_n^{(m)} & \text{if } y_m(x_n) = t_n \\ w_n^{(m)} \exp \alpha_m & \text{Otherwise} \end{cases}$$

3. Classify the datapoints as

$$Y_M(\mathbf{x}) = \text{sign} \sum_{m=1}^M \alpha_m y_m(\mathbf{x})$$

Adaptive boosting: the algorithm

1. Set $w_n^{(1)} = \frac{1}{N}$
2. For $m = 1, \dots, M$
 - 2.1 Fit $y_m(\mathbf{x})$ by minimising $E_m = \sum_{n \in \mathcal{M}_m} w_n^{(m)}$

2.2 Evaluate

$$\epsilon_m = \frac{\sum_{n \in \mathcal{M}_m} w_n^{(m)}}{\sum_n w_n^{(m)}},$$


set $\alpha_m = \log \frac{1-\epsilon_m}{\epsilon_m}$

2.3 Update the weights

$$w_n^{(m+1)} = \begin{cases} w_n^{(m)} & \text{if } y_m(x_n) = t_n \\ w_n^{(m)} \exp \alpha_m & \text{Otherwise} \end{cases}$$

3. Classify the datapoints as

$$Y_M(\mathbf{x}) = \text{sign} \sum_{m=1}^M \alpha_m y_m(\mathbf{x})$$



Adaptive boosting: the algorithm

1. Set $w_n^{(1)} = \frac{1}{N}$
2. For $m = 1, \dots, M$
 - 2.1 Fit $y_m(\mathbf{x})$ by minimising $E_m = \sum_{n \in \mathcal{M}_m} w_n^{(m)}$

2.2 Evaluate

$$\epsilon_m = \frac{\sum_{n \in \mathcal{M}_m} w_n^{(m)}}{\sum_n w_n^{(m)}},$$


$$\text{set } \alpha_m = \log \frac{1 - \epsilon_m}{\epsilon_m}$$

2.3 Update the weights


$$w_n^{(m+1)} = \begin{cases} w_n^{(m)} & \text{if } y_m(x_n) = t_n \\ w_n^{(m)} \exp \alpha_m & \text{Otherwise} \end{cases}$$


3. Classify the datapoints as


$$Y_M(\mathbf{x}) = \text{sign} \sum_{m=1}^M \alpha_m y_m(\mathbf{x})$$





- ▶ Simple 2-class, 2-dimensional dataset
- ▶ Simple classifiers:
 - ▶ Choose threshold θ , dimension d , class c
 - ▶ $y(\mathbf{x}) = \begin{cases} \mathcal{C}_c & \text{iff } x_d < \theta \\ \mathcal{C}_{1-c} & \text{Otherwise} \end{cases}$
- ▶ Create ensemble using adaboost



















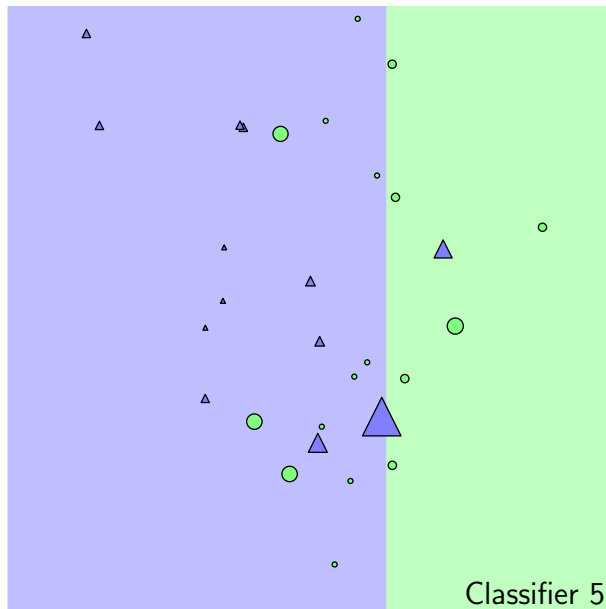




























Classifier 7










Adaboost can be interpreted as minimising

$$E = \sum_{n=1}^N \exp \left(-\frac{t_n}{2} \sum_{m=1}^M \alpha_m y_m(\mathbf{x}_n) \right)$$


As a consequence:

1. It strongly penalises misclassifications, not robust to outliers!
2. It does not generalise to more than 2 classes
3. Choosing a different error function
 - ▶ Allows multiclass classification and even regression (e.g. Gradient Boosting)
 - ▶ Makes robust classifiers possible




A nice application of boosting:

- ▶ Very simple features (HAAR wavelets)
 - ▶ Use integral images to compute these very fast
- ▶ Use *cascading* for speedup




A nice application of boosting:

- ▶ Very simple features (HAAR wavelets)
 - ▶ Use integral images to compute these very fast
- ▶ Use *cascading* for speedup



A nice application of boosting:

- ▶ Very simple features (HAAR wavelets)
 - ▶ Use integral images to compute these very fast
- ▶ Use *cascading* for speedup



Introduction

Bias-Variance Decomposition

Bagging and Boosting


Bagging


Boosting


Tree-based models

Classification trees

Random Forests







Tree-based models split the input space in regions

- ▶ Each region gets its own classifier
- ▶ The classifiers can be extremely simple (typically: constant)


Pros	Cons
<ul style="list-style-type: none">▶ Interpretable!▶ Simple and fast▶ If let to grow, will learn perfect classification on the training data▶ Pruning (using validation set) allows proper generalisation	<ul style="list-style-type: none">▶ Final tree depends strongly on particular data▶ Hard decisions, aligned with dimensions▶ Finding best tree is intractable

Combine trees with bagging and random feature selection


Procedure: for N datapoints and M features,
pre-specify $m \ll M$

1. Repeat K times:
 - 1.1 Get a bootstrap sample
 - 1.2 At each node in the tree:
 - 1.2.1 select m features at random
 - 1.2.2 Find the optimal split based on these m features and the training set
 - 1.3 Fully grow the tree (no pruning)

This is often considered one of the most powerful committee methods


$$M = 1$$


$M = 2$




$M = 2$

$M = 10$




$M = 10$

$M = 50$




$M = 100$



$M = 100$

$M = 100$ Final decision



Introduction

Bias-Variance Decomposition

Bagging and Boosting

Bagging

Boosting

Tree-based models

Classification trees

Random Forests

To summarise:

- ▶ Combine models to improve their expressive power (cfr. Mixture of Gaussians)
- ▶ Combining independent models can dramatically improve performance
- ▶ Making different models responsible for different areas of the space combines simple models into very flexible models